



REACTIVE « JAMMED » ARCHITECTURE

ou comment survivre à l'A86 en heure de pointe



Nicolas PHUNG  @nsphung



Jennifer AGUIAR  @yaourtcorp

The COSMO project



Next generation e-checking system

- Nationwide train passenger e-checking
- Near real time
- Offline device capabilities
- Work with complex existing services

Crowded Services/Partners

Crowded services/partners

- Galaxies of services
- Each services got its speed limit and various QoS
- We are not the only customer !
- Some services are already jammed with lots of layers



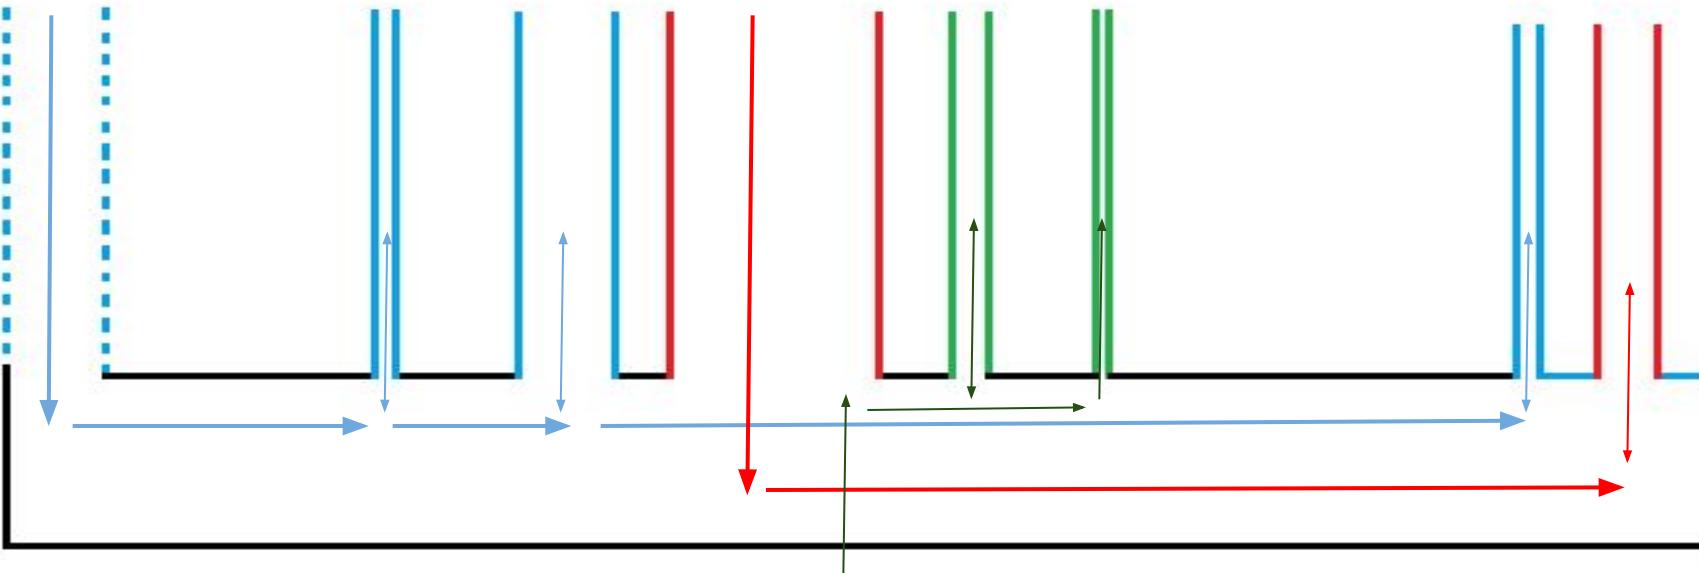
Different Constraints

- “Real Time”
- Scalable (sometimes)
- Resilient
- Batch !?

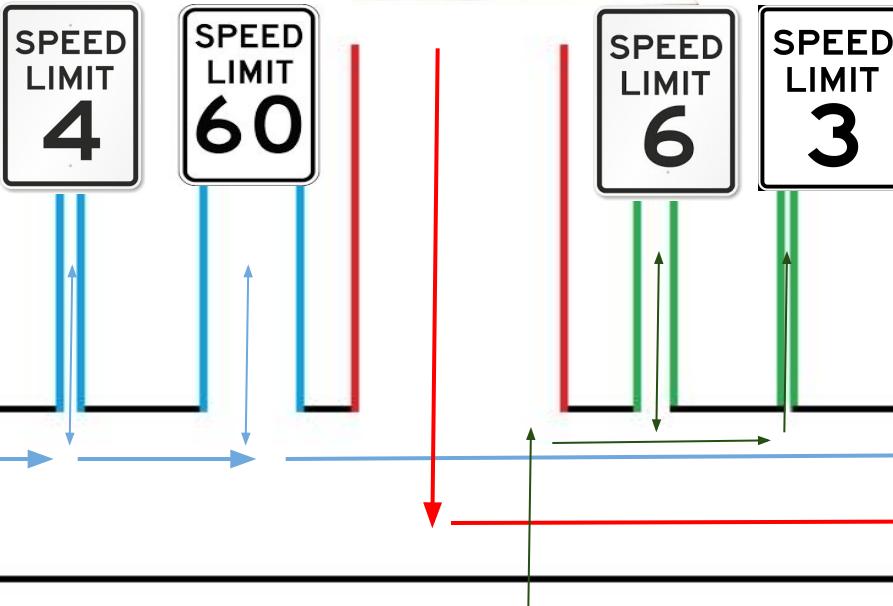
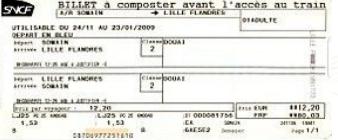


The A-team | mertcaliskan | CC3

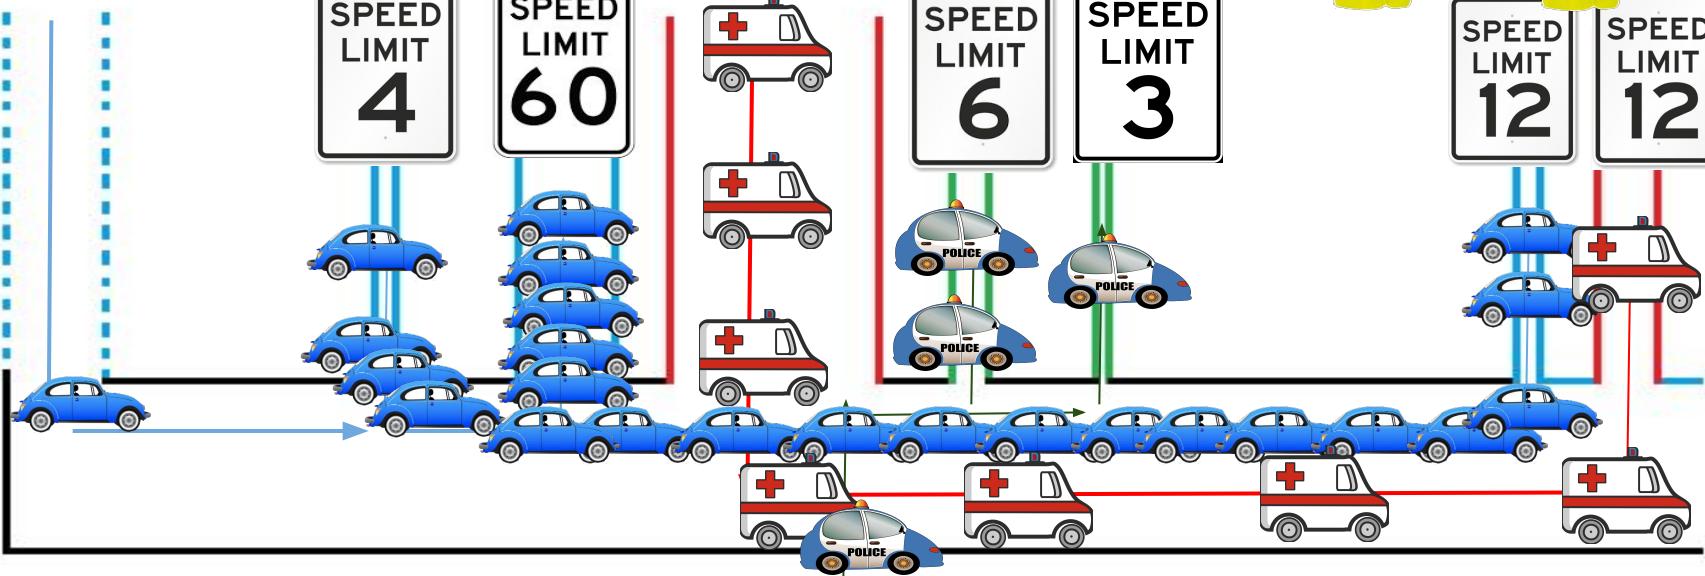
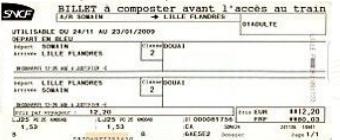
ROADS...MAP



+ SPEED...LIMITS



= TRAFFIC...JAM



REQUIREMENTS

PL	<3,5t	>3,5t
	50 5:00-23:00 60 23:00-5:00	50 5:00-23:00 60 23:00-5:00
	90	70
	100	80
	120	* 100
	140	

KAFKA and AKKA to the rescue



CONCEPTS THAT WILL HELP US

- Speed limit => Throttle ✓
- Keep up with jammed services => Backpressure
- Address near real time concern => Event based Architecture

Back to the pressure

Dès que le trafic devient plus dense, la circulation ralentit.

Dès que le trafic devient plus dense, la circulation ralentit. Jusqu'à l'arrêt ou presque. Chacun s'attend à apercevoir l'accident ou les travaux à l'origine du ralentissement quand, sans raison apparente, la situation se débloque. Chacun reprend alors son rythme de croisière. Mais quelques kilomètres plus loin, ça recommence... "Les 'accordéons' se forment quand la circulation est déjà congestionnée, explique Christine Buisson, du Laboratoire d'ingénierie Circulation Transports (Licit), à Vaulx-en-Velin. Dans ces conditions, la vitesse d'un véhicule est imposée par le véhicule qui le précède. Si elle ralentit, les autres aussi."

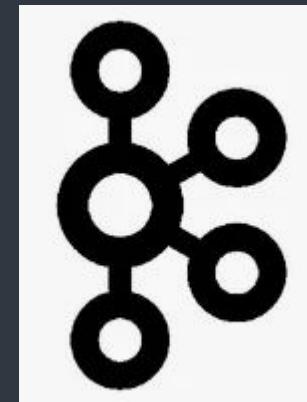
L'explication vient de la dynamique des fluides, car les voitures en circulation sont comme un fluide de n'importe quelle matière qui s'écoule dans un contenant, un tuyau par exemple, à la capacité limitée. Ainsi, les théoriciens du trafic routier utilisent les lois de l'hydrodynamique pour simuler la circulation des véhicules.

une vague de compression qui se propage dans la circulation vers l'arrière !!!

"D'un point de vue macroscopique, on voit bien dans nos simulations que si le trafic passe rapidement de 2 000 à 4 000 véhicules par heure, une compression se forme qui se propage comme une vague." **Une vague qui présente la particularité de déferler dans le sens inverse de la circulation!**

Back to the pressure (Part II)

We may need a big buffer...

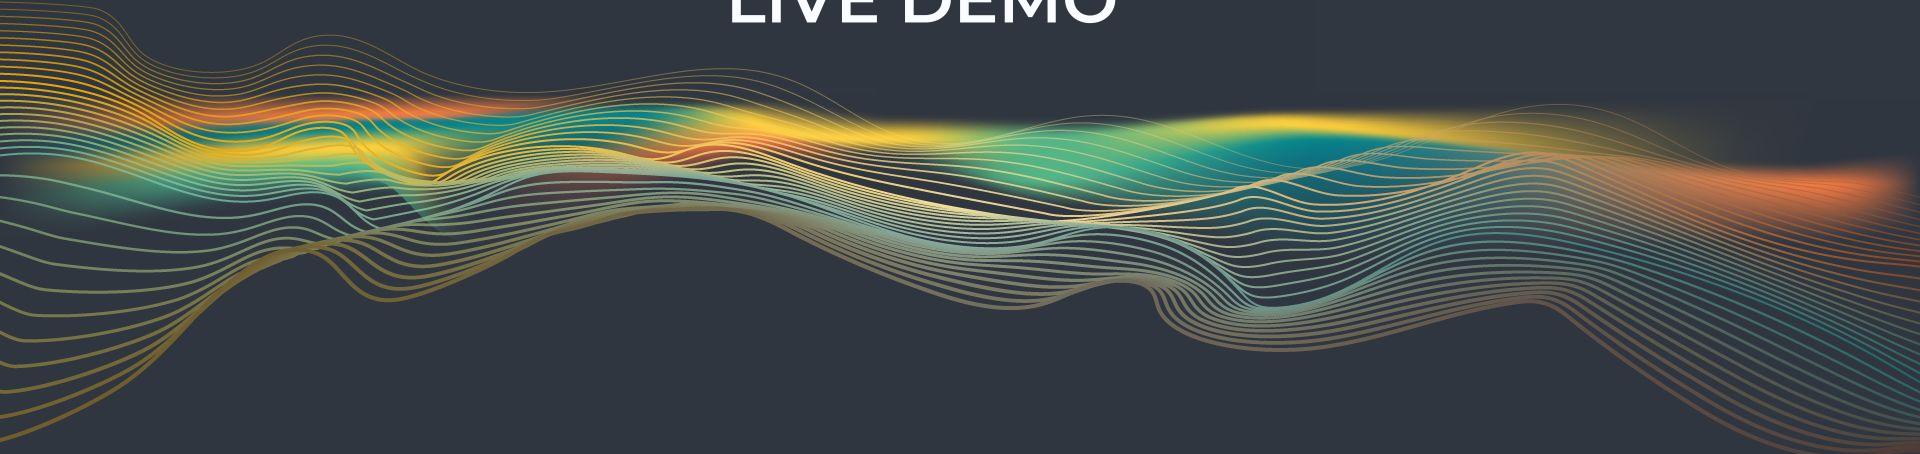


Insane Traffic jam | Splloid

Back to the pressure (Part III)

Is this a holdup?
No, it's a science experiment!

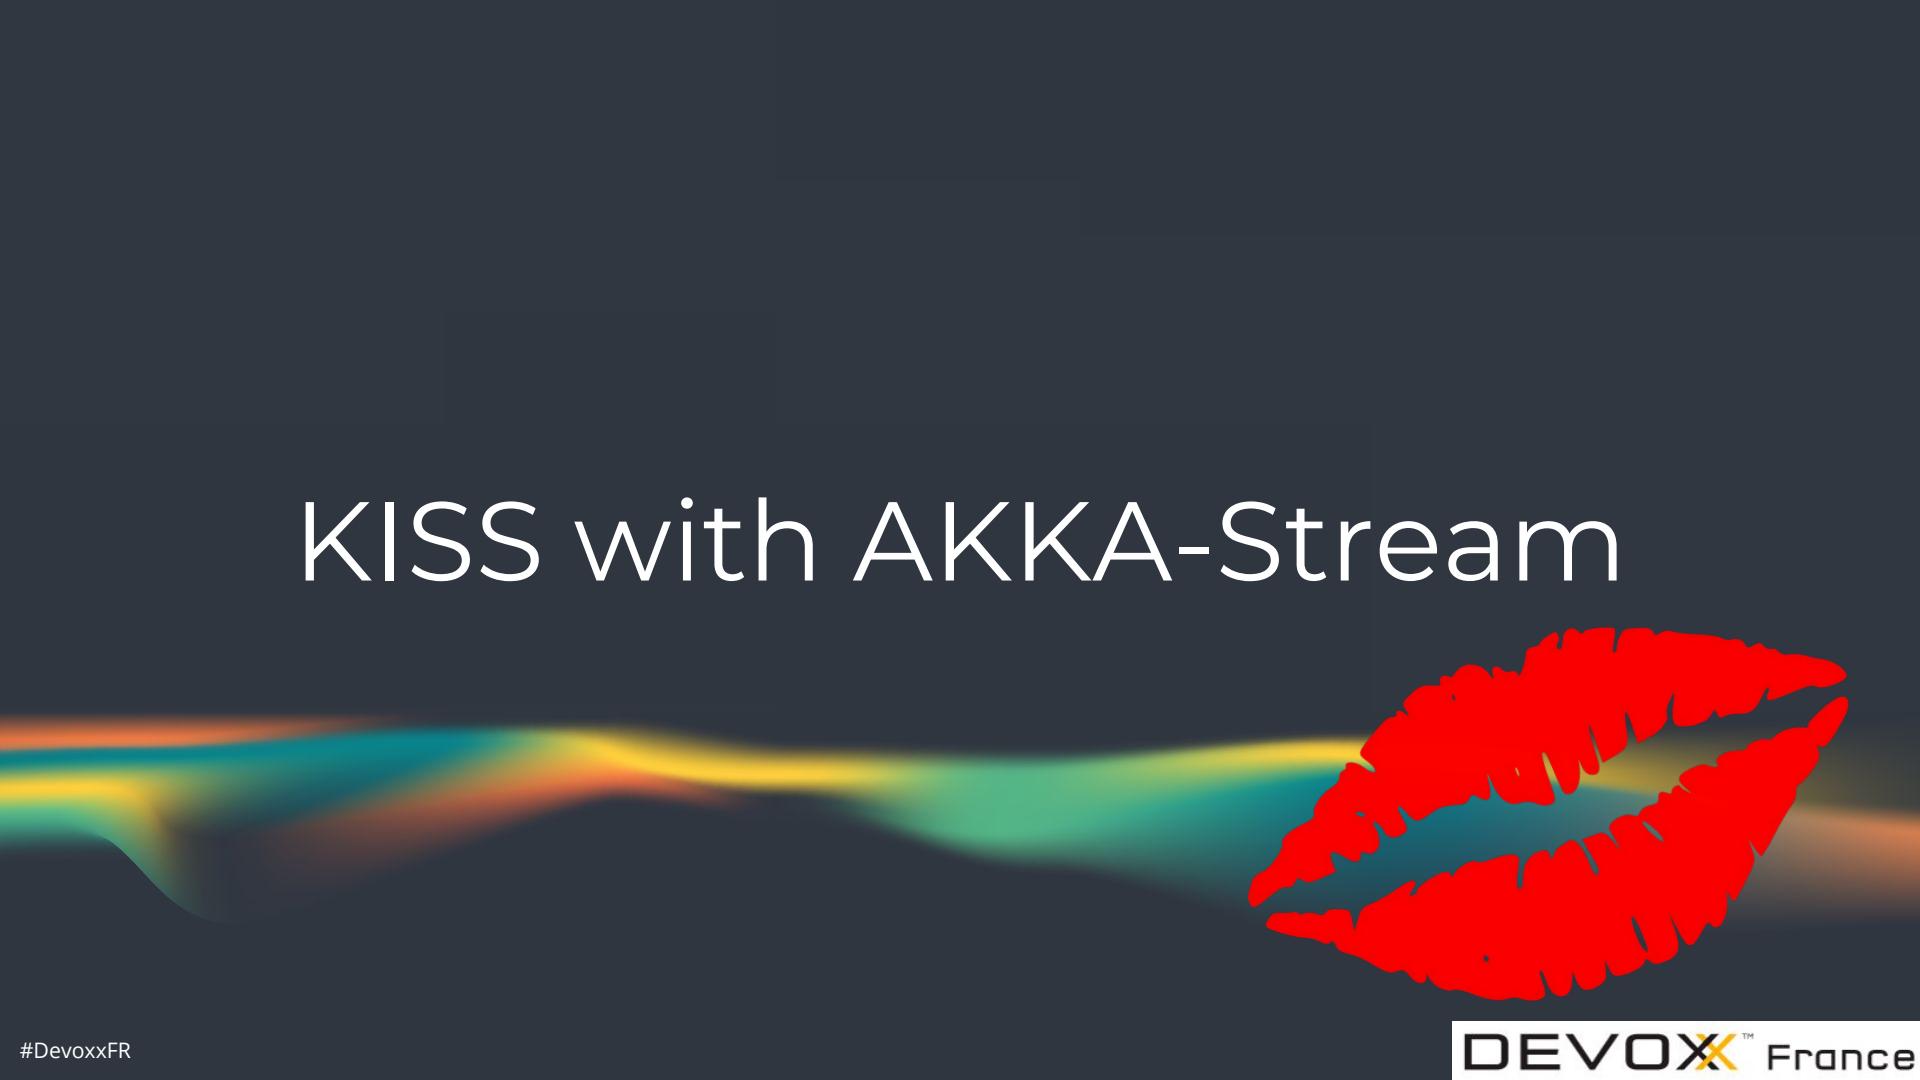
LIVE DEMO



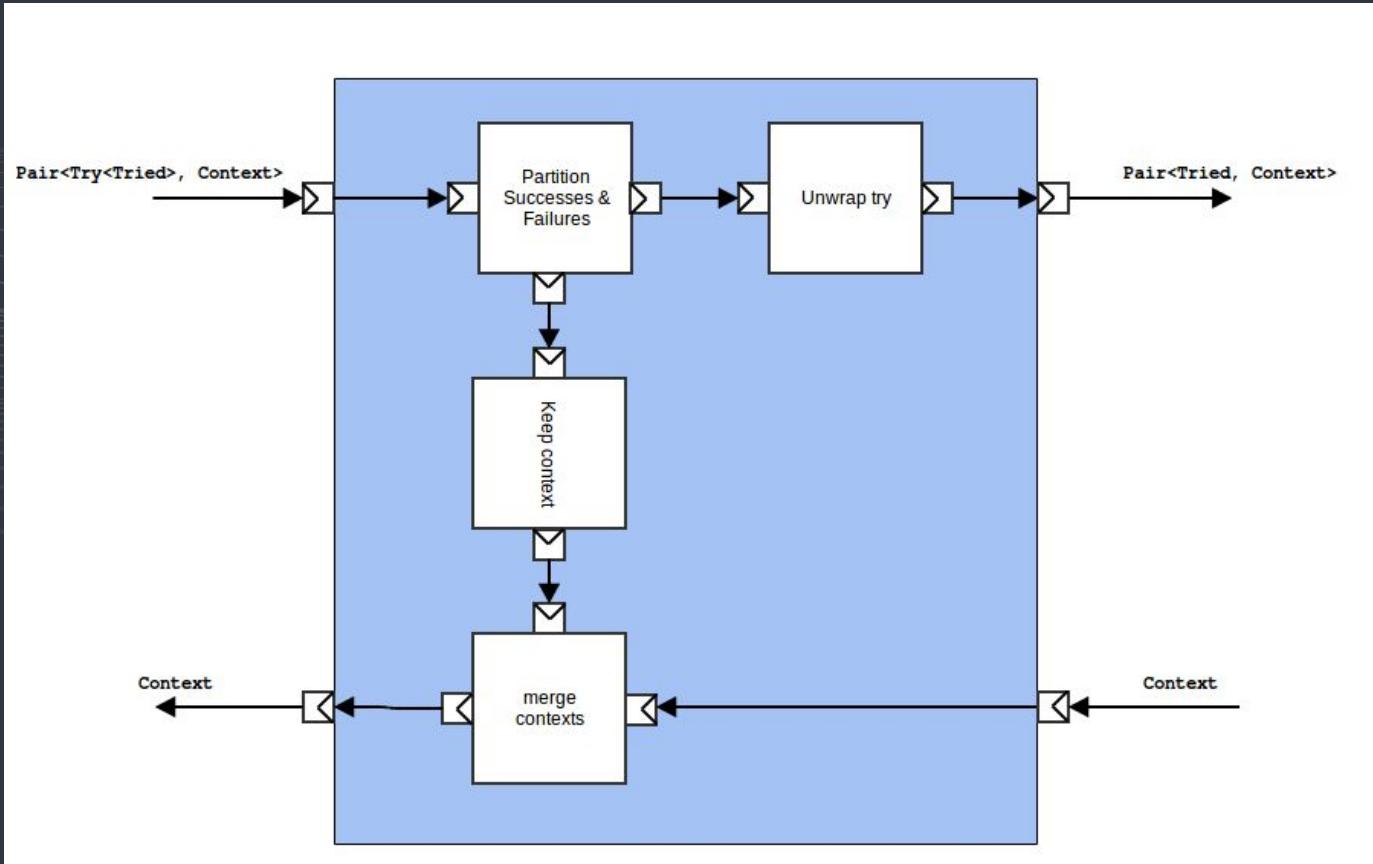
CONCEPTS THAT WILL HELP US

- Speed limit => Throttle ✓
- Keep up with jammed services => Backpressure ✓
- Address near real time concern => Event-based Architecture

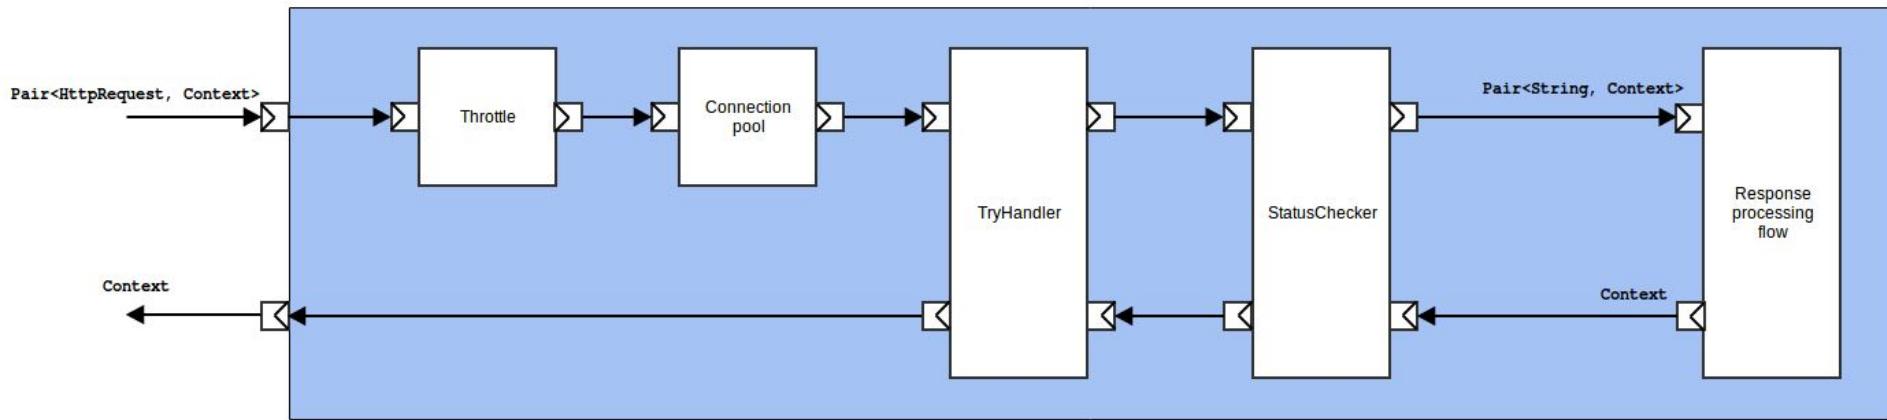
KISS with AKKA-Stream



FIRST VERSION

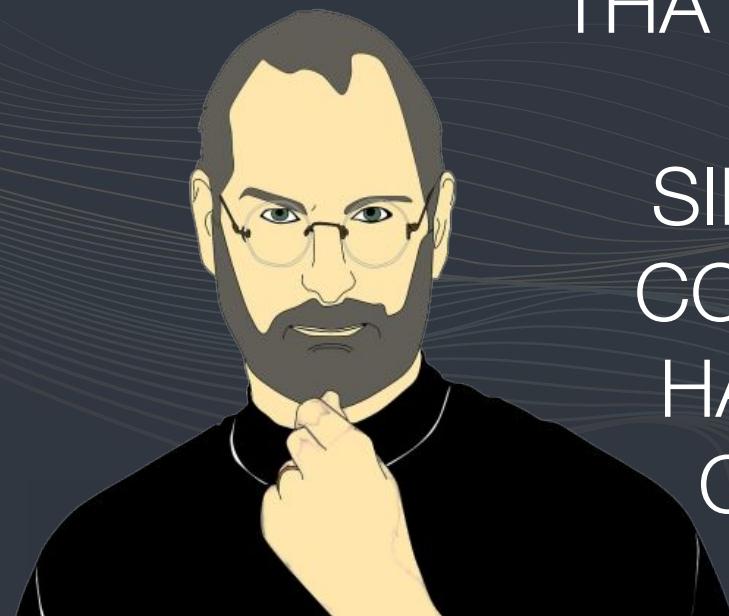


Not so V



Lots of code / clutter components

Really difficult to test/evolve

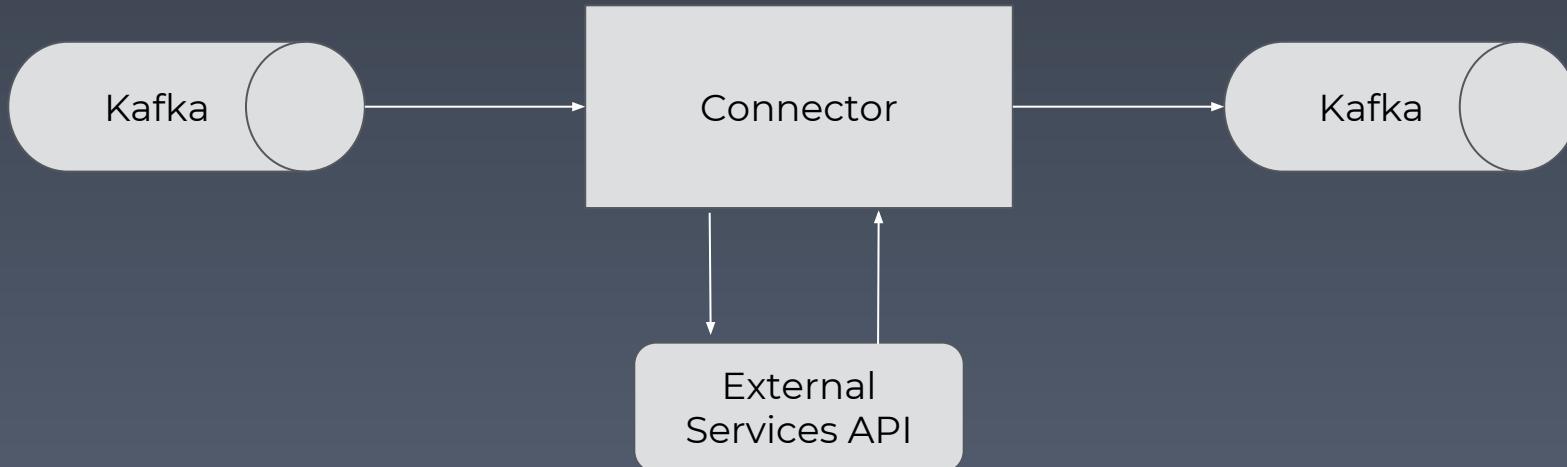


THAT'S BEEN ONE OF MY MANTRAS
FOCUS & SIMPLICITY

SIMPLE CAN BE HARDER THAN
COMPLEX; YOU HAVE TO WORK
HARD TO GET YOUR THINKING
CLEAN AND MAKE IT SIMPLE

- STEVE JOBS

Main use case



Simpler version

```
val graph = source
  .collect(CollectSpecificTypeInEnvelope[OrderRequestEvent])
  .filter(keepEventAfterYesterdaySessionId(_, MARKER, skipPastDays))
  .via(buildStreamContextFlow)
  .throttle(throttling.maxRequestsPerSecond, throttling.period, throttling.maximumBurst, throttling.throttleMode)
  .mapAsync(mapAsyncParallelism)(CallOrderDetailsPaoWebService(orderDetailsUtils, redisService, paoSettings))
  .mapConcat[(OrderDetailsContext, Option[TicketScope]]](pairEventsToContextWithOptionalOffsetTupled)
  .via(buildKafkaProducerMessageFlow)
  .via(producerFlow)
  .map(elem => elem.message.passThrough)
  // -- Extract only OrderByTrainContext with Some committable offset
  .collect { case OrderDetailsContext(_, _, _, _, Some(offset), _) => offset }
  .via(commitOffsetBackToKafkaTopicSourceFlow(commitParallelism))
  .withAttributes(ActorAttributes.withSupervisionStrategy(streamDecider))
  .to(sink)
```



Simpler version

Read from kafka

```
val graph = source
    .collect(CollectSpecificTypeInEnvelope[OrderRequestEvent])
    .filter(keepEventAfterYesterdaySessionId(_, MARKER, skipPastDays))
    .via(buildStreamContextFlow)
    .throttle(throttling.maxRequestsPerSecond, throttling.period, throttling.maximumBurst, throttling.throttleMode)
    .mapAsync(mapAsyncParallelism)(CallOrderDetailsPaoWebService(orderDetailsUtils, redisService, paoSettings))
    .mapConcat[(OrderDetailsContext, Option[TicketScope]]](pairEventsToContextWithOptionalOffsetTupled)
    .via(buildKafkaProducerMessageFlow)
    .via(producerFlow)
    .map(elem => elem.message.passThrough)
    // -- Extract only OrderByTrainContext with Some committable offset
    .collect { case OrderDetailsContext(_, _, _, _, Some(offset), _) => offset }
    .via(commitOffsetBackToKafkaTopicSourceFlow(commitParallelism))
    .withAttributes(ActorAttributes.withSupervisionStrategy(streamDecider))
    .to(sink)
```

Simpler version

```
val graph = source
  .collect(CollectorsType)
  .filter(keepEventAfterYesterday)
  .via(buildStreamContextFlow)
  .throttle(throttling.maxRequestsPerSecond, throttling.period, throttling.maximumBurst, throttling.throttleMode)
  .mapAsync(mapAsyncParallelism)(CallOrderDetailsPaoWebService(orderDetailsUtils, redisService, paoSettings))
  .mapConcat[(OrderDetailsContext, Option[TicketScope])](pairEventsToContextWithOptionalOffsetTupled)
  .via(buildKafkaProducerMessageFlow)
  .via(producerFlow)
  .map(elem => elem.message.passThrough)
  // -- Extract only OrderByTrainContext with Some committable offset
  .collect { case OrderDetailsContext(_, _, _, _, Some(offset), _) => offset }
  .via(commitOffsetBackToKafkaTopicSourceFlow(commitParallelism))
  .withAttributes(ActorAttributes.withSupervisionStrategy(streamDecider))
  .to(sink)
```

Filter stuff

Simpler version

```
val graph = source
  .collect(CollectSpecificTypeInEnvelope[OrderRequestEvent])
  .filter(keepEventAfterYesterday)
  .via(buildStreamContextFlow) Add some context ( correlationId, offset... )
  .throttle(throttling.maxRequestsPerSecond, throttling.period, throttling.maximumBurst, throttling.throttleMode)
  .mapAsync(mapAsyncParallelism)(CallOrderDetailsPaoWebService(orderDetailsUtils, redisService, paoSettings))
  .mapConcat[(OrderDetailsContext, Option[TicketScope]]](pairEventsToContextWithOptionalOffsetTupled)
  .via(buildKafkaProducerMessageFlow)
  .via(producerFlow)
  .map(elem => elem.message.passThrough)
  // -- Extract only OrderByTrainContext with Some committable offset
  .collect { case OrderDetailsContext(_, _, _, _, Some(offset), _) => offset }
  .via(commitOffsetBackToKafkaTopicSourceFlow(commitParallelism))
  .withAttributes(ActorAttributes.withSupervisionStrategy(streamDecider))
  .to(sink)
```

Simpler version

```
val graph = source
  .collect(CollectSpecificTypeInEnvelope[OrderRequestEvent])
  .filter(keepEventAfterYesterdaySessionId(_, MARKER, skipPastDays))
  .via(buildStreamContextFlow)
  .throttle(throttling.maxRequest // Speeeeeed limit!
    .mapAsAsync(mapAsAsyncParallelism)(CallOrderDetailsPaoWebService(orderDetailsUtils, redisService, paoSettings))
    .mapConcat[(OrderDetailsContext, Option[TicketScope]]](pairEventsToContextWithOptionalOffsetTupled)
    .via(buildKafkaProducerMessageFlow)
    .via(producerFlow)
    .map(elem => elem.message.passThrough)
    // -- Extract only OrderByTrainContext with Some committable offset
    .collect { case OrderDetailsContext(_, _, _, _, Some(offset), _) => offset }
    .via(commitOffsetBackToKafkaTopicSourceFlow(commitParallelism))
    .withAttributes(ActorAttributes.withSupervisionStrategy(streamDecider))
  .to(sink)
```

Simpler version

```
val graph = source
  .collect(CollectSpecificTypeInEnvelope[OrderRequestEvent])
  .filter(keepEventAfterYesterdaySessionId(_, MARKER, skipPastDays))
  .via(buildStreamContextFlow)
  .throttle(throttling.maxRequestPerSecond * 1000)
  .mapAsync(mapAsyncParallelism)
  .mapConcat[(OrderDetailsContext, Option[TicketScope]])(pairEventsToContextWithOptionalOffsetTupled)
  .via(buildKafkaProducerMessageFlow)
  .via(producerFlow)
  .map(elem => elem.message.passThrough)
  // -- Extract only OrderByTrainContext with Some committable offset
  .collect { case OrderDetailsContext(_, _, _, _, Some(offset), _) => offset }
  .via(commitOffsetBackToKafkaTopicSourceFlow(commitParallelism))
  .withAttributes(ActorAttributes.withSupervisionStrategy(streamDecider))
  .to(sink)
```

Call webservice in // but don't forget the order please!

Simpler version

```
val graph = source
  .collect(CollectSpecificTypeInEnvelope[OrderRequestEvent])
  .filter(keepEventAfterYesterdaySessionId(_, MARKER, skipPastDays))
  .via(buildStreamContextFlow)
  .throttle(throttling.maxRequestsPerSecond, throttling.period, throttling.maximumBurst, throttling.throttleMode)
  .mapAsync(mapAsyncParallelism)(callOrderDetailsRawWebService(orderDetailsService, redisService, nanoSettings))
  .mapConcat[(OrderDetailsContext] Ok, this one's tricky! We'll talk about sashimis just after
  .via(buildKafkaProducerMessageFlow)
  .via(producerFlow)
  .map(elem => elem.message.passThrough)
  // -- Extract only OrderByTrainContext with Some committable offset
  .collect { case OrderDetailsContext(_, _, _, _, Some(offset), _) => offset }
  .via(commitOffsetBackToKafkaTopicSourceFlow(commitParallelism))
  .withAttributes(ActorAttributes.withSupervisionStrategy(streamDecider))
  .to(sink)
```

Simpler version

```
val graph = source
  .collect(CollectSpecificTypeInEnvelope[OrderRequestEvent])
  .filter(keepEventAfterYesterdaySessionId(_, MARKER, skipPastDays))
  .via(buildStreamContextFlow)
  .throttle(throttling.maxRequestsPerSecond, throttling.period, throttling.maximumBurst, throttling.throttleMode)
  .mapAsync(mapAsyncParallelism)(CallOrderDetailsPaoWebService(orderDetailsUtils, redisService, paoSettings))
  .mapConcat[(OrderDetailsContext, Option[TicketScope])](pairEventsToContextWithOptionalOffsetTupled)
  .via(buildKafkaProducerFlow)
  .via(producerFlow) Write to kafka
  .map(elem => elem.message.passThrough)
  // -- Extract only OrderByTrainContext with Some committable offset
  .collect { case OrderDetailsContext(_, _, _, _, Some(offset), _) => offset }
  .via(commitOffsetBackToKafkaTopicSourceFlow(commitParallelism))
  .withAttributes(ActorAttributes.withSupervisionStrategy(streamDecider))
  .to(sink)
```

Simpler version

```
val graph = source
  .collect(CollectSpecificTypeInEnvelope[OrderRequestEvent])
  .filter(keepEventAfterYesterdaySessionId(_, MARKER, skipPastDays))
  .via(buildStreamContextFlow)
  .throttle(throttling.maxRequestsPerSecond, throttling.period, throttling.maximumBurst, throttling.throttleMode)
  .mapAsync(mapAsyncParallelism)(CallOrderDetailsPaoWebService(orderDetailsUtils, redisService, paoSettings))
  .mapConcat[(OrderDetailsContext, Option[TicketScope]]](pairEventsToContextWithOptionalOffsetTupled)
  .via(buildKafkaProducerMessageFlow)
  .via(producerFlow)
  .map(elem => elem.message.passThrough
    // -- Extract only OrderByTrainContent
    .collect { case OrderDetailsContext(
      .via(commitOffsetBackToKafkaTopicSource)
      .withAttributes(ActorAttributes.withSupervisionStrategy(streamDecider))
      .to(sink)
```

Just commit the kafka offset if you can but in
batching mode!

Simpler version

```
val graph = source
  .collect(CollectSpecificTypeInEnvelope[OrderRequestEvent])
  .filter(keepEventAfterYesterdaySessionId(_, MARKER, skipPastDays))
  .via(buildStreamContextFlow)
  .throttle(throttling.maxRequestsPerSecond, throttling.period, throttling.maximumBurst, throttling.throttleMode)
  .mapAsync(mapAsyncParallelism)(CallOrderDetailsPaoWebService(orderDetailsUtils, redisService, paoSettings))
  .mapConcat[(OrderDetailsContext, Option[TicketScope]]](pairEventsToContextWithOptionalOffsetTupled)
  .via(buildKafkaProducerMessageFlow)
  .via(producerFlow)
  .map(elem => elem.message.passThrough)
  // -- Extract only OrderByTrainContext with Some committable offset
  .collect { case OrderDetailsContext(_, _, _, _, Some(offset), _) => offset }
  .via(commitOffsetBackToKafkaTopicSourceFlow(commitParallelism))
  .withAttributes(ActorAttributes.withSupervision)
  .to(sink)
```

Don't forget the “in case of emergency” part...

Simpler version

```
val graph = source
  .collect(CollectSpecificTypeInEnvelope[OrderRequestEvent])
  .filter(keepEventAfterYesterdaySessionId(_, MARKER, skipPastDays))
  .via(buildStreamContextFlow)
  .throttle(throttling.maxRequestsPerSecond, throttling.period, throttling.maximumBurst, throttling.throttleMode)
  .mapAsync(mapAsyncParallelism)(CallOrderDetailsPaoWebService(orderDetailsUtils, redisService, paoSettings))
  .mapConcat[(OrderDetailsContext, Option[TicketScope]]](pairEventsToContextWithOptionalOffsetTupled)
  .via(buildKafkaProducerMessageFlow)
  .via(producerFlow)
  .map(elem => elem.message.passThrough)
  // -- Extract only OrderByTrainContext with Some committable offset
  .collect { case OrderDetailsContext(_, _, _, _, Some(offset), _) => offset }
  .via(commitOffsetBackToKafkaTopicSourceFlow(commitParallelism))
  .withAttributes(ActorAttributes.withSupervisionStrategy(streamDecider))
  .to(sink)
```

... and go nowhere!

Just linear graph logic is
enough !

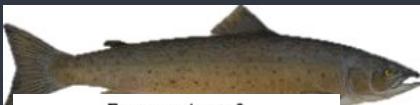
Get salmons from the river



Frozen-salmon-1



Powered By IDAutomation.com



Frozen-salmon-2



Powered By IDAutomation.com



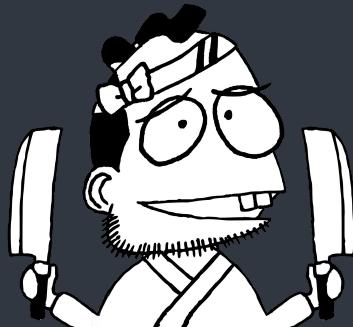
Frozen-salmon-3



Powered By IDAutomation.com



My Chef / Webservices



Tracing SASHIMIS...



Frozen-salmon-1



Powered By IDAutomation.com



Frozen-salmon-2



Powered By IDAutomation.com



Frozen-salmon-3



Powered By IDAutomation.com



Powered By IDAutomation.com



Frozen-salmon-2

Frozen-salmon-2



Powered By IDAutomation.com



Frozen-salmon-3

Frozen-salmon-3



Frozen-salmon-1



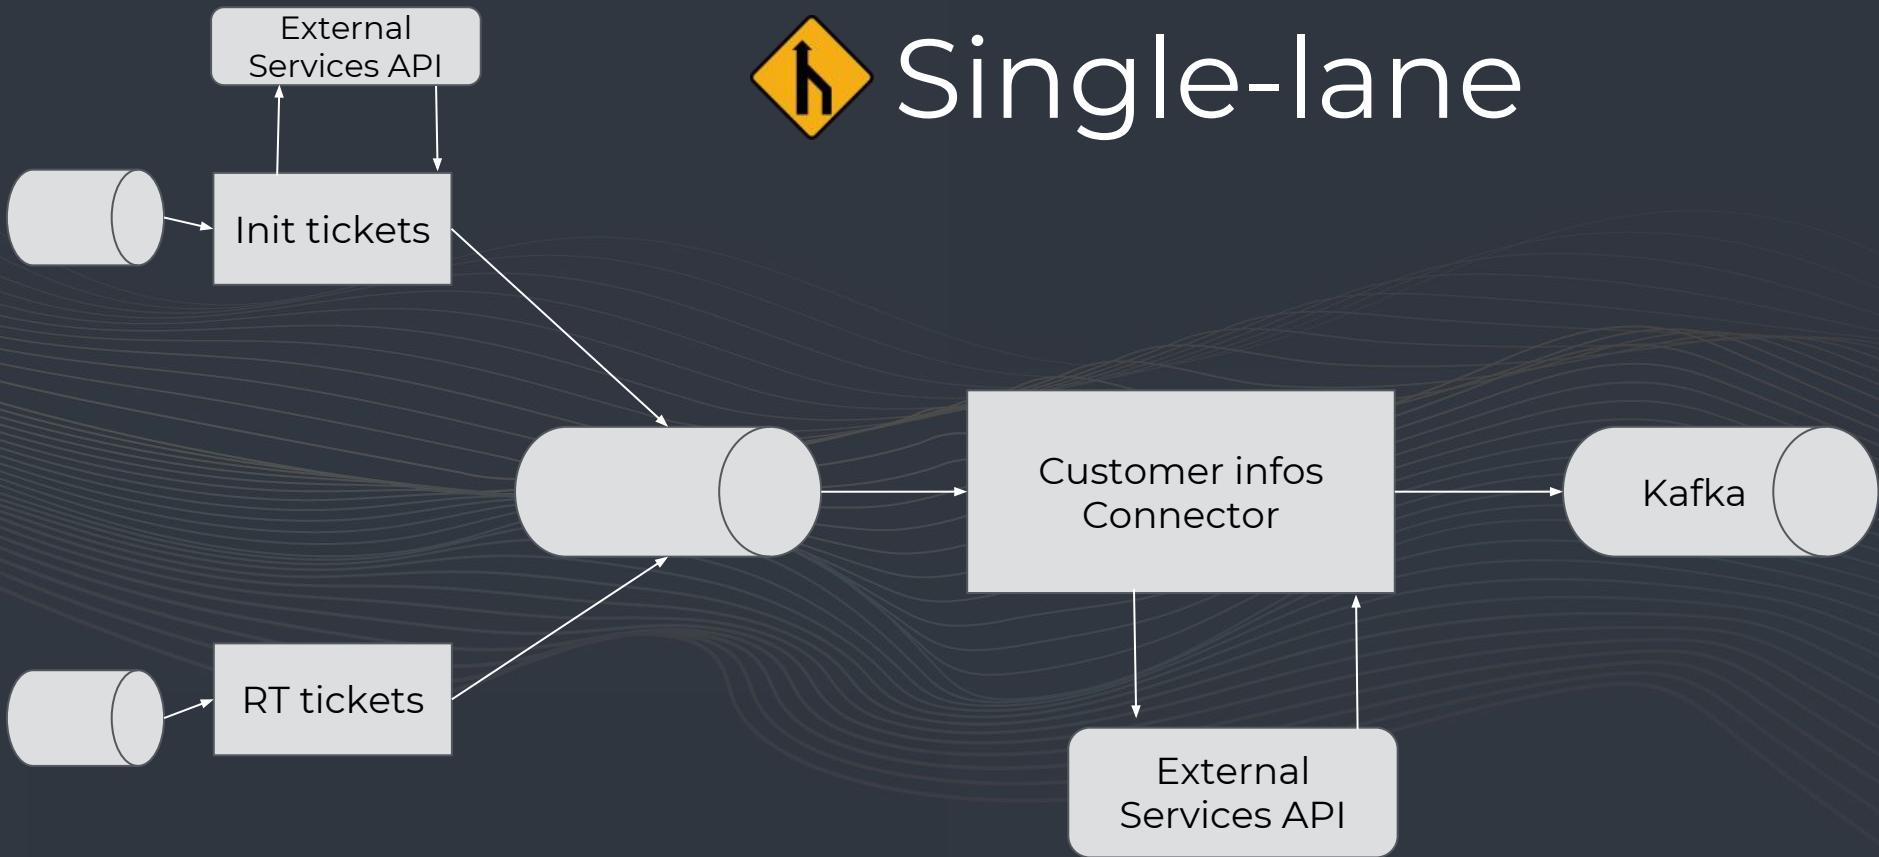
Powered By IDAutomation.com

Speed up “real time”





Single-lane



Rush hour

Initialization processing sends
many many many more events
than the so-called “Real Time” one

“Real time” events are lost in the
transportation flow



Image by [Golda Falk](#) from [Pixabay](#)

Rush hour

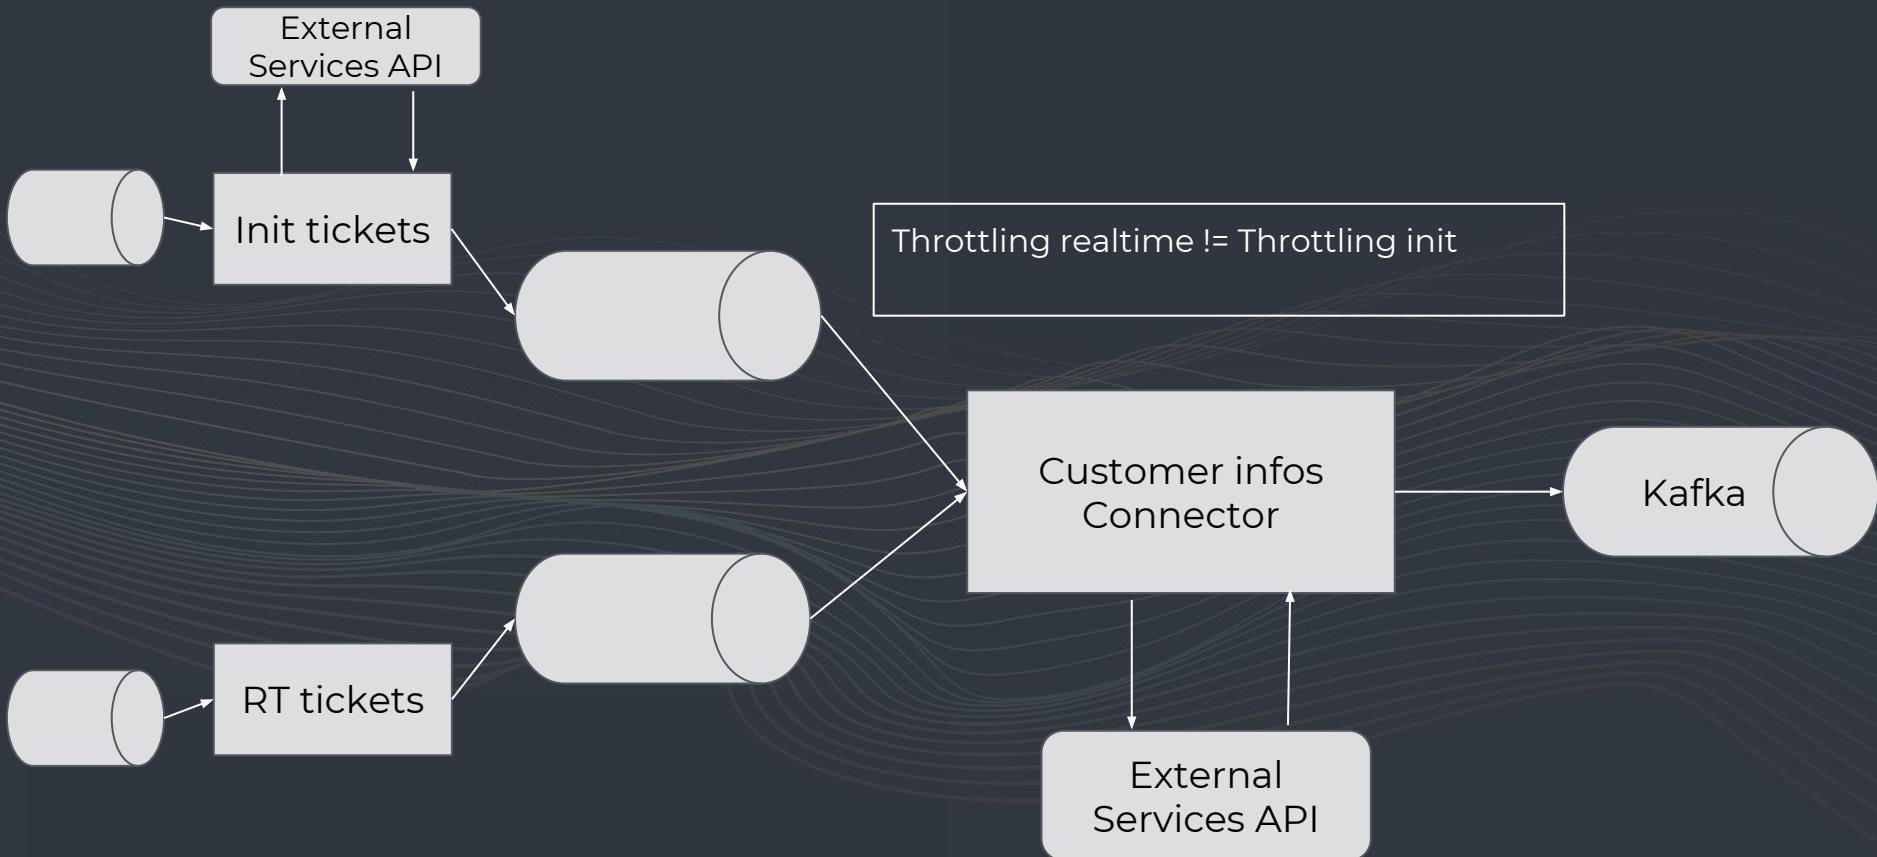
Initialization processing sends
many many many more events
than the so-called “Real Time” one

“Real time” events are lost in the
transportation flow

And these critical events are late...
really late...



Still waiting for my train | Henry Burrows | CC2



So basically, we've just
“created” a bus lane...



Into the wild code

```
trait EmeraudeDiscovery extends ServiceDiscovery {
    def configuration: Configuration

    def redisService: RedisService

    override def zookeeperConnect: String = configuration.getZookeeperConnectionString

    override def partnerName: String = Partner.EMR.name()

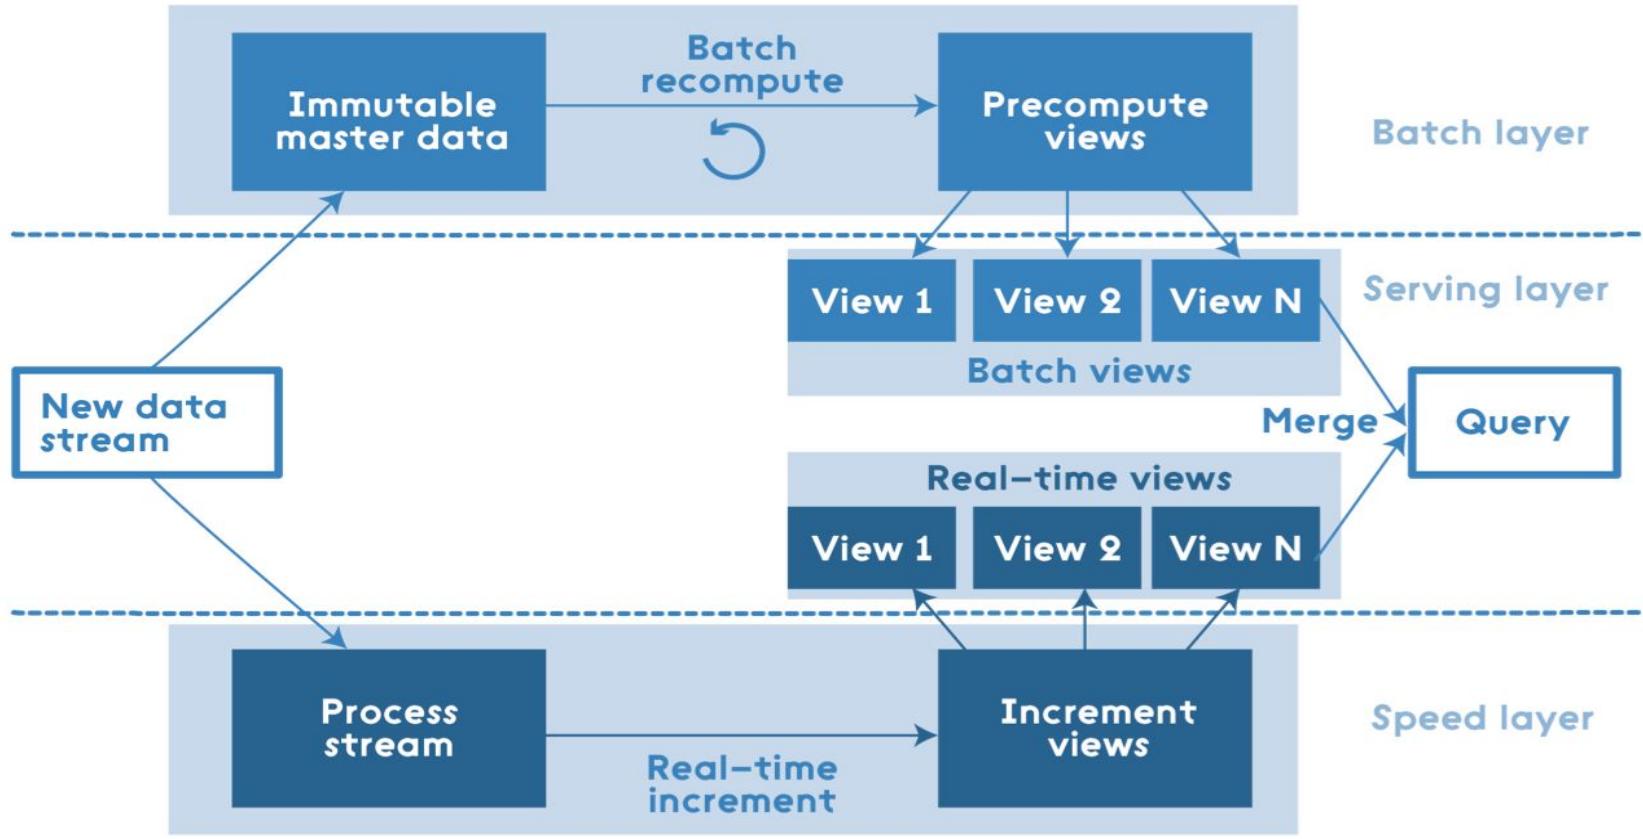
    override def serverPort: Int = configuration.getServerPort

    override def createStreams(): Seq[StreamDiscovery with StreamControl] = {
        // -- Register your Stream here
        val passengersInfoDispatchStreamGraph = new PassengersInfoDispatchStreamGraph()
        val passengerDataInitStreamGraph = new PassengerDataStreamGraph("INIT", redisService, "streams.passenger-data-init")
        val passengerDataRealTimeStreamGraph = new PassengerDataStreamGraph("REAL_TIME", redisService, "streams.passenger-data-real-time")

        Seq(
            passengersInfoDispatchStreamGraph,
            passengerDataInitStreamGraph,
            passengerDataRealTimeStreamGraph
        )
    }
}
```

It's the same code but with a different configuration! (throttling, topics...)

Look like Lambda Architecture

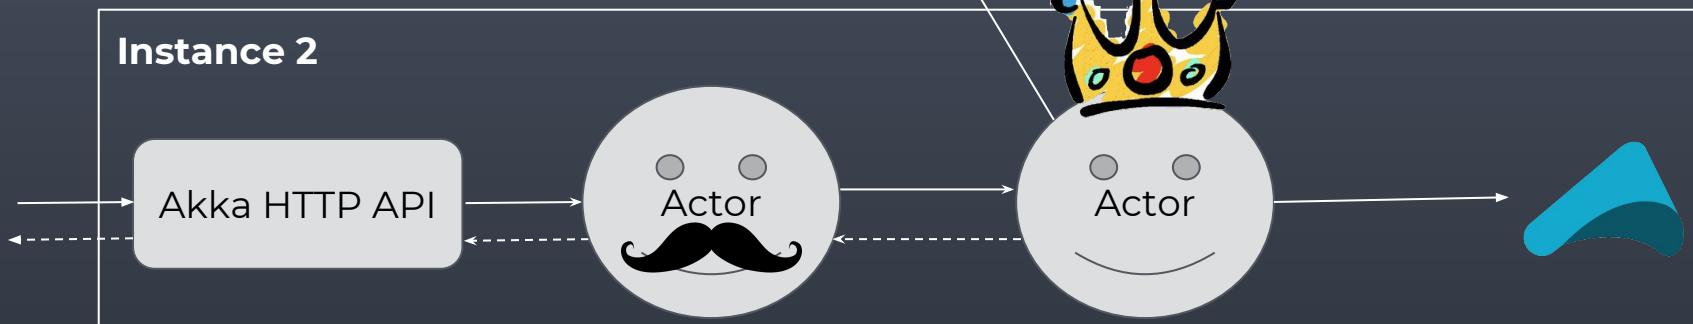
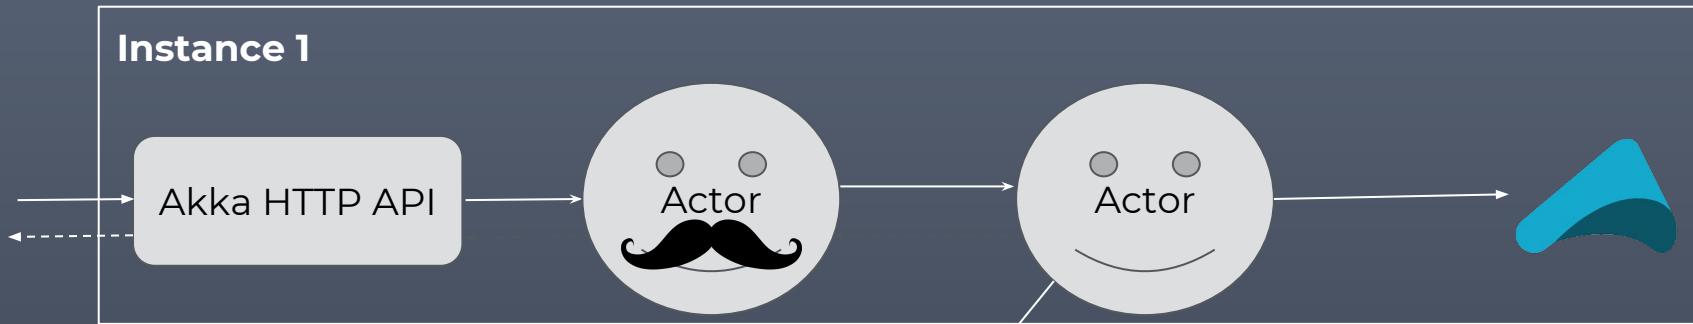


Operating streams

[Photo by Jonathan Wheeler on Unsplash](#)

Operating streams

- Throttling is static in akka-stream : when your graph is built and running, you cannot change it dynamically!
- Sometimes it may start, stop, restart...
- Or the stream is a batch





Feedback

LEARNING

- Try to keep simple graph (as easy to read and maintain)
- Use Akka DSL only when necessary
- Avoid clutter components when simple programming language is enough

LEARNING

- Use metadata on your message to measure, identify, correlate...
- Measure, measure, measure!
- Beware of lag

LEARNING

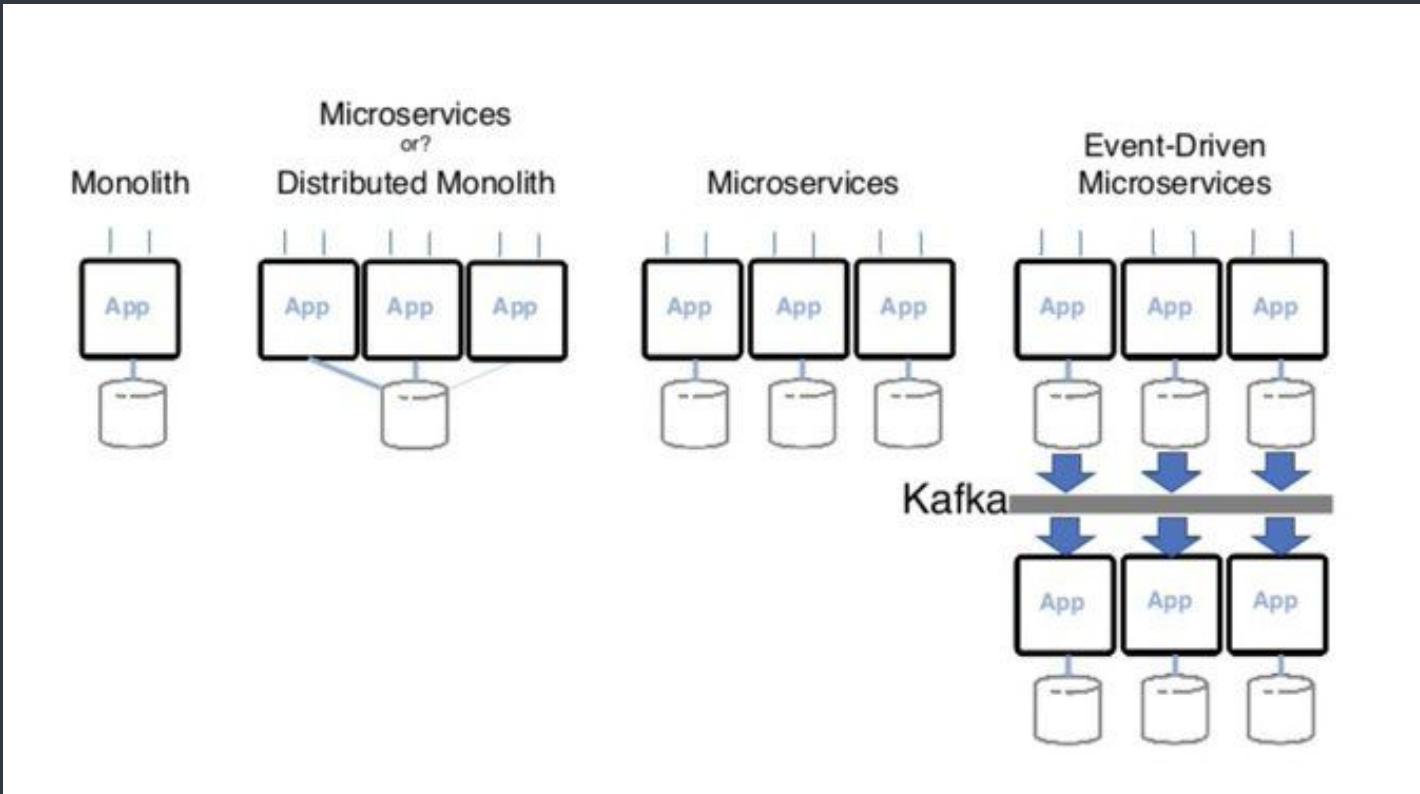
Sometimes backpressure is your friend and sometimes it's not...

LEARNING

Sometimes backpressure is your friend and sometimes it's not...



Architecture Evolution



Source: <https://twitter.com/cdavisafc/status/1105572564712210432>

CONCEPTS THAT WILL HELP US

- Speed limit => Throttle ✓
- Keep up with jammed services => Backpressure ✓
- Address near real time concern => Event-based Architecture ✓

Thank you!



Questions ?

ANNEXES



**Backpressure and cache
adventure with COSMO Backend**

<https://nsphung.netlify.com/csm-backpressure/>

“AKKA FTW” -Jonas Boner

Akka Team
@akkateam
Akka Team @ Lightbend
akka.io
Joined December 2011

Tweets **Tweets & replies** **Media**

↳ Akka Team Retweeted
Jonas Bonér @jboner · 11h
#Akka FTW.
"...will power tomorrow unified control device on all SNCF Trains!"

Nicolas Phung @nsphung
Just presented Reactive "Jammed" Architecture with @yaourtcorp on last friday at internal Oui Tech Conference @ouisncf_news Trying to be The A-Team / "L'agence tout risque" with akka-streams and Kafka. This #backend...

Akka* scales?



Lightbend  @lightbend

Following 

How to scale? Here you go:

With over 125 million players, and supporting over 8 million *concurrent gaming sessions*, we are really happy to learn that #Fortnite is running #Akka under the hood!

Dragos Manolescu @hysteresis

Epic's #fortnite running on AWS and "powered by Java, #Akka and even Go"
#reInvent

9:24 AM - 27 Nov 2018

70 Retweets 151 Likes



 1

 70

 151



58

DEVOXX™ France

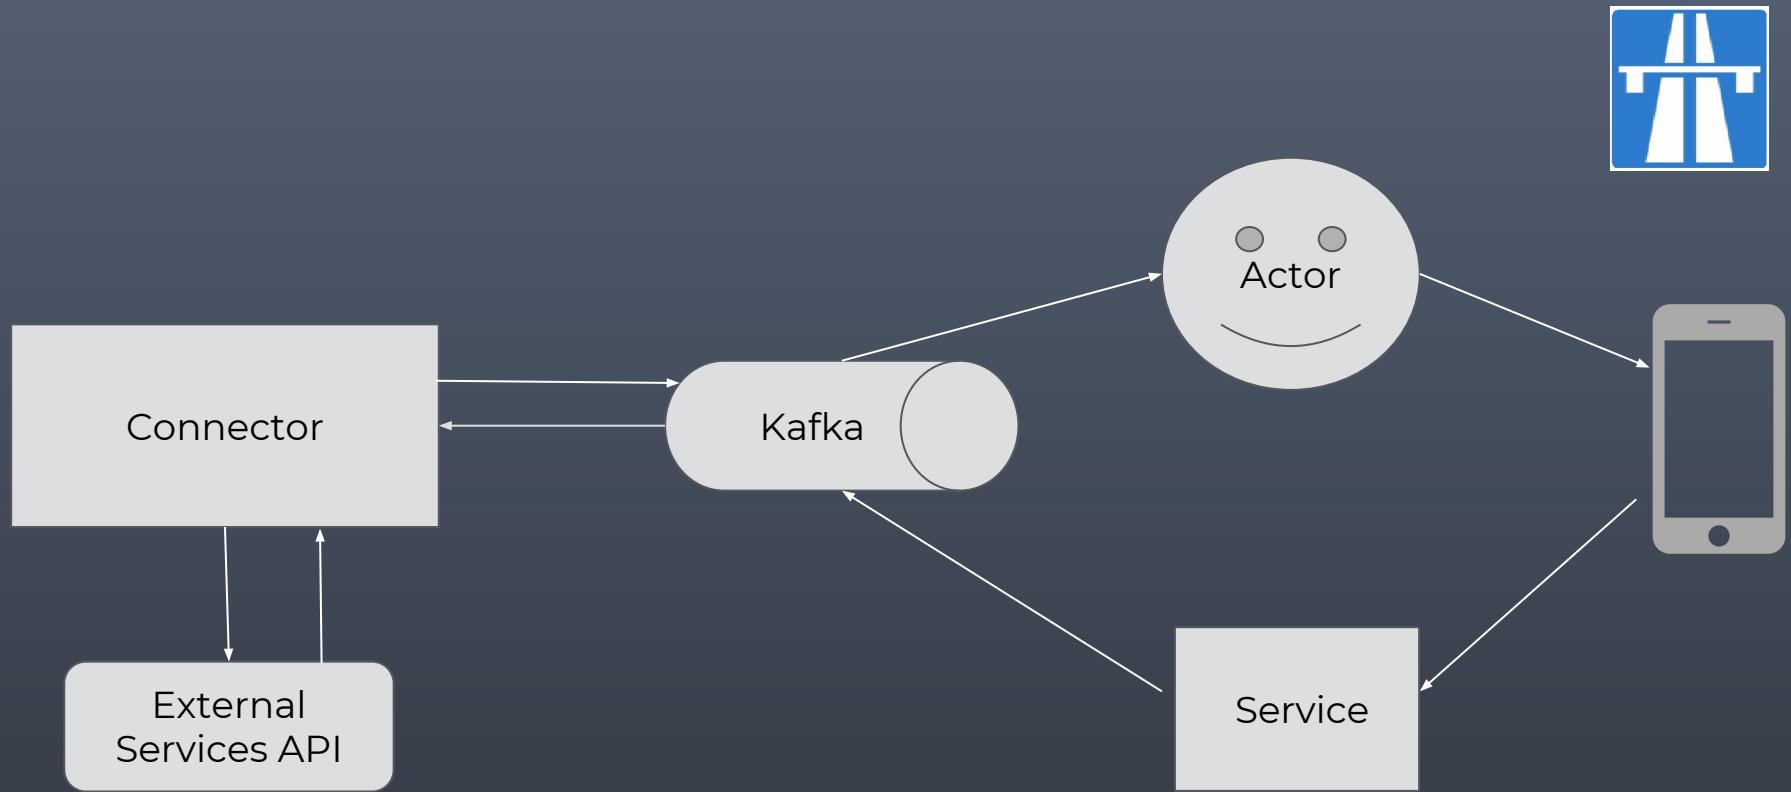
#DevoxxFR

Checkmate with Actors



Tickets Check

- Handle events from partners and expose them to a mobile front-end
- Receive events from a mobile front-end and forward them to a partner
- All these events may affect the same data
- Everything has an end!



LEARNING



- Separate your reads and your writes
- Sometimes RAM is enough (and sometimes not!)
- Your partners may fail but you should keep the “closest to reality” internal state
- Actors are fine for keeping “no-lifetime” data (kafka is here for that) but think about boot time and possible inconsistencies between instances